

Concurrent Model Serving



Diagram illustrating Concurrent Model Serving:

```


graph TD
    A[ollama] --> B["load [Model A]"]
    A --> C["load [Model B]"]
    B --> D["embedding"]
    C --> E["embedding"]
    D --> F["Memory"]
    E --> G["Memory"]
    F --> H["Inference"]
    G --> I["Inference"]
    H --> J["Memory"]
    I --> K["Memory"]
    J --> L["Embedding"]
    K --> M["Embedding"]
    L --> N["User 1"]
    M --> O["User 2"]
  
```

This architecture allows the server to receive and process inference requests for any of these loaded models without incurring the "cold start" latency of loading and unloading from disk for each request.

A server that is not concurrent would face this "cold start" penalty every time it needed to switch tasks. This makes any professional application unusably slow.



Model Chaining




The Benefits: Why This Is a Requirement, Not a "Nice-to-Have"

You don't "chain" just because you can; you do because it's the only way to solve certain problems.

Task Specialization + Higher Quality

Model chaining enables each model to focus on its area of expertise, leading to higher-quality results. Instead of forcing one model to perform multiple, conflicting tasks, each model in the chain handles a specific part of the workflow it excels at. This avoids performance trade-offs and ensures that every step — whether technical, analytical, or creative — is completed with maximum precision and depth. The result is a combined output that's more accurate, coherent, and polished than what any single model could produce alone.


Specialized



Enabling Agents Behaving (The "Critic" Pattern)
Model chaining is often used in systems where the system acts like an autonomous agent rather than a passive responder. Instead of producing a single static output, models can engage in reflection and self-improvement loops. In this pattern — often referred to as the "Critic" loop — a model generates a critique based on previous interactions, requests feedback improvements, and then revises its output based on this feedback. This iterative process leads to refined and refined, leading to more refined, accurate, and consistent work.

Such feedback-driven workflows form the foundation of advanced multi-agent systems like **Heteron**, where models continuously collaborate and evolve toward optimal solutions.

Agentic Behavior → models behaving like an **autonomous agent** → Human Being



Cost Optimization through Selective Model Usage
Cost Optimization through Selective Model Usage involves strategically leveraging models of varying capabilities throughout a workflow. For simpler or routine tasks — such as initial data extraction, basic classification, or filtering — you can use lighter, faster, and more cost-effective models. The more powerful, resource-intensive models are then reserved for **final stages**. This approach not only reduces overall computational costs but also ensures that high-end resources are allocated efficiently, maximizing both speed and performance across the workflow.

→ **Selectively** → Model (Small, Medium, Big).

Improved Debugging and Traceability

Chained workflows significantly enhance debugging and traceability. When an unexpected result occurs, it becomes much easier to pinpoint the source of the problem. By examining the inputs and outputs of each individual step in the workflow, you can identify exactly where an error or misstep happened. This is far more efficient than attempting to interpret the internal state of a single, large model call, allowing for faster troubleshooting and more precise corrections.

